KUNAL KHADSE

Sr Data Engineer | Microsoft

Senior Data Engineering professional with 10+ years of experience designing, modernizing, and governing enterprise-scale Azure and Databricks platforms. Expertise spans Lakehouse Architecture, Data Governance, Data Integration, Cloud Modernization, and Analytics Engineering. Proven track record advising strategic enterprise customers, driving architecture decisions, and leading complex cloud transformation initiatives. Passionate about enabling organizations to build scalable, secure, and data-driven platforms.

.

kkhadse31@gmail.com | +48 739 627 018 | Poland

PERSONAL INTEREST

Kunal enjoys adventure trekking and exploring new places.

EDUCATION & TRAINING

2011 – 2015

Bachelor of Engineering (Computer Science)

CERTIFICATIONS

DP-203: Data Engineering on Microsoft Azure

DP-400 / DP-204 — Azure DevOps Expert

AI-901

DP-750 -Data Engineering solution with Azure databricks

AZ-900: Azure Fundamentals

Databricks: Unified Analytics Platform

Azure AI Fundamentals

PROFESSIONAL SUMMARY

Senior Data Engineer with 10+ years of experience architecting and scaling enterprise data platforms on Azure and Databricks, now working at the intersection of technical leadership and solutions architecture. Proven track record leading complex, cross-functional escalations for global enterprise customers, driving platform adoption, and translating business requirements into scalable technical architecture. Deep expertise across the Modern Data Platform stack — Databricks, Spark, Delta Lake, Unity Catalog, Azure Data Factory — combined with strong stakeholder engagement, mentoring, and Agile delivery leadership. Seeking a Lead Data Engineer / Solutions Architect role to drive technical strategy, architecture decisions, and platform excellence at scale.

SKILLS SUMMARY

Cloud & Data Platforms	Data Engineering	Programming	DevOps & Automation	Architecture & Governance	Analytics & AI
Azure	PySpark	Python	Azure DevOps	Power BI	Lakehouse Architecture
Azure Databricks	Spark SQL	SQL	Gitlab	Azure AI Foundry	Data Governance
Azure Synapse	Delta Lake		Git	OpenAI Integrations	Data Security
Azure Data Factory	Unity Catalog		Docker	AI	Data Lineage
Azure Event Hub	DLT, Genie, Lakebase		Kubernetes		Copilot

WORK EXPERIENCE

Senior Data Engineer — Microsoft *Jan 2023 – Present*

- Act as a trusted technical advisor for strategic enterprise customers, driving architecture decisions and modernization initiatives across Azure Databricks data platforms.
- Architect and implement scalable Lakehouse solutions using Azure Databricks, Unity Catalog, Delta Lake, and Azure analytics services.
- Define enterprise-wide governance frameworks encompassing data security, lineage, access management, and compliance controls.
- Lead architecture reviews and solution design workshops, translating business requirements into scalable cloud data platform architectures.
- Partner with Product Engineering, Supportability, and Customer Engineering teams to resolve critical platform challenges and improve service reliability.
- Establish reusable architecture patterns and best practices for data ingestion, transformation, governance, and analytics workloads.
- Built an AI-driven knowledge management platform using GitHub Copilot and Azure AI technologies, improving solution discovery and reducing engineering effort across globally distributed teams.

- Influence strategic customer adoption of emerging Azure Databricks capabilities, including Unity Catalog, Genie, Lakebase, and AI-powered analytics solutions.

Technologies: Azure Databricks, Unity Catalog, Delta Lake, Lakehouse Architecture, Azure Synapse, Azure AI Foundry, GitHub Copilot, OpenAI, PySpark, Data Governance, Solution Architecture

Senior Data Engineer — Luxoft

Feb 2022 – Jan 2023

- Led the modernization of P&G's enterprise data warehouse by designing and implementing a cloud-native Azure-based Modern Data Platform supporting analytics and reporting workloads.
- Architected and delivered a centralized data hub integrating multiple business domains, enabling self-service analytics and reducing dependency on legacy reporting systems.
- Built distributed data processing frameworks using PySpark and Spark SQL to handle large-scale structured and semi-structured datasets.
- Implemented Medallion Architecture (Bronze, Silver, Gold layers) to improve data quality, governance, and downstream analytics performance.
- Collaborated with business stakeholders, data analysts, and reporting teams to translate business requirements into scalable and reusable data products.
- Optimized data pipelines and transformation workloads, reducing overall processing times and improving platform reliability.
- Established data quality controls, monitoring mechanisms, and operational best practices to improve trust and consistency across reporting systems.
- Contributed to cloud migration and modernization initiatives, helping transition legacy warehouse workloads to Azure-based analytics platforms.

Technologies: Azure Databricks, Azure Data Factory, PySpark, Spark SQL, Delta Lake, Azure Data Lake Storage Gen2, Power BI, Python, SQL, Git, CI/CD, Modern Data Warehouse, Medallion Architecture

Technologies: Databricks, Azure Data Factory, GIT, Azure DevOps

Lead data Engineer — Accenture

May 2020 – Feb 2022

- Built a centralized data hub for a banking client to consolidate and visualize enterprise data via Power BI.
- Designed Azure Data Factory pipelines and led ingestion and Databricks-based transformation workflows across multiple linked services.
- Containerized and deployed applications using Docker, publishing artifacts to a Nexus registry, and configured an AKS cluster with a full CI/CD deployment pipeline.
- Analyzed and transformed large-scale Parquet datasets using PySpark on Databricks.

Technologies: Docker, Kubernetes, Databricks, Azure Data Factory, GIT, Azure DevOps, TeamCity

Senior Software Engineer — Yash Technologies

Apr 2019 – Apr 2020

- Delivered a centralized data hub and Power BI reporting layer for a UK-based client.
- Designed Azure Data Factory pipelines and led ingestion and Databricks transformation activity across linked services.
- Built a Delta Lake layer to support upsert/merge patterns for incremental data processing.
- Implemented real-time streaming ingestion of sensor data using Event Hub and IoT Hub.

Technologies: Spark, Python, Databricks, Azure Data Factory, GIT, Azure DevOps

Big Data Developer — TCS

Oct 2018 – Apr 2019

- Assessed data integrity for a US-based manufacturing client's Nexia server using PySpark, and visualized findings in Tableau.
- Built Azure Data Factory pipelines and parsed and transformed Nexia server data into Parquet-format tables using PySpark on Databricks.
- Integrated supply-unit and parts-unit data across IWD, R12, P21, and EMIA data warehouses, and led migration of Oracle Apps invoice and order data into the Enterprise Data Hub.
- Used Spark SQL to transform, clean, and filter imported data, loading curated tables into the CNF zone in Parquet format.
- Automated table refresh logic via Unix shell scripting to identify updated partitions and drive targeted drop/insert operations; scheduled jobs using Alteryx.

Technologies: Spark SQL, Hive, Sqoop, Alteryx, JIRA

Hadoop Developer — TCS

Oct 2015 – Apr 2018

- Gathered requirements from business stakeholders and translated them into technical specifications, working closely with developers to validate functional design.

- Imported data from Oracle R12 using Sqoop and built Hive-based analysis layers with Parquet storage, including external/internal tables and Hive optimization techniques.
- Performed transformation, cleaning, and filtering of imported data using Hive and MapReduce, loading curated output into HDFS.

Technologies: Impala, Hive, Sqoop, JIRA

ACHIEVEMENTS

- Customer Champion Award, 2024–2025 H1 - Microsoft
- Out of the box Award - Microsoft
- Extra Miler achievement - Accenture
- Star Performer Award from client -Yash technology
- NVIDIA-approved CUDA Developer